Evidence Gate - SHOP-102 (updated requirement) - commit b7d0762 · Raise free-shipping threshold

BLOCKED

3 MUST criteria not verified

ROUTE

back to developer

AC-2, AC-3, AC-4 failed

Next action: Fix the code so these hold, then rerun evidence-gate check: AC-2 "An order of $74.99 pays standard shipping ($5.99)." (shipping: expected 5.99, got 0.0); AC-3 "An order of $50.00 no longer ships free (old rule removed)." (shipping: expected 5.99, got 0.0); AC-4 "The customer's receipt shows the correct shipping on both sides of $75." (shipping: expected 5.99, got 0.0).

> Warning: The working tree has uncommitted changes; the evidence is for the files as they are now.

Criteria

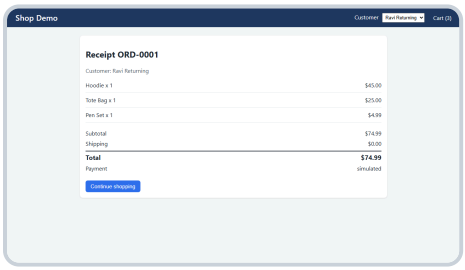
CRITERION	PRIORITY	LABEL	EVIDENCE
AC-1	MUST	Needs human review	subtotal=75.00 -> shipping 0.00 (expected 0.00); How should subtotals be rounded or compared when they have more than two decimals, such as 74.995? The spec does not define this.
AC-2	MUST	Not verified: failed	subtotal=74.99 -> shipping 0.00 (expected 5.99)
AC-3	MUST	Not verified: failed	subtotal=50.00 -> shipping 0.00 (expected 5.99); Should any order between $50.00 and $74.99 ever ship free, for example through a legacy promotion or for particular customer types? The spec says the old rule is removed, but the diff adds a legacy $50 branch.
AC-4	MUST	Not verified: failed	customer=returning-customer-01, add=HOODIE-001,BAG-001,PEN-001 -> shipping 0.00 (expected 5.99); customer=returning-customer-01, add=HOODIE-001,BAG-001,NOTE-001 -> shipping 0.00 (expected 0.00); Do all customer types, such as new, guest and returning customers, follow the same threshold? AC-4 only covers returning-customer-01.
(new)		Needs human review	Is the threshold checked against the subtotal before or after discounts, taxes or promo codes? The spec does not say.
(new)		Needs human review	What is the expected shipping for a zero or negative subtotal, such as an empty cart or a refund? The spec does not say.
(new)		Needs human review	Does the quote API (AC-1 to AC-3) use the same shipping function as the browser receipt path (AC-4)? Is any other threshold copy, such as in the UI or docs, expected to change too?

Evidence: every scenario that was run

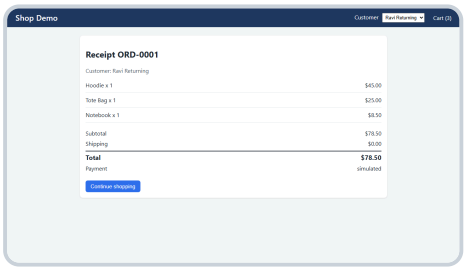
Only scenarios with an expected value written by a person count towards a label.

SCENARIO	INPUT	EXPECTED	ACTUAL	STATUS
S-AC-1-1 AC-1 · MUST	subtotal=75.00	shipping 0.00	shipping 0.00	PASS
S-AC-2-1 AC-2 · MUST	subtotal=74.99	shipping 5.99	**shipping 0.00 ≠ 5.99**	FAIL
S-AC-3-1 AC-3 · MUST	subtotal=50.00	shipping 5.99	**shipping 0.00 ≠ 5.99**	FAIL
S-AC-4-1 AC-4 · MUST	customer=returning-customer-01, add=HOODIE-001,BAG-001,PEN-001	shipping 5.99	**shipping 0.00 ≠ 5.99**	FAIL
S-AC-4-2 AC-4 · MUST	customer=returning-customer-01, add=HOODIE-001,BAG-001,NOTE-001	shipping 0.00	shipping 0.00	PASS

Screenshots



S-AC-4-1



S-AC-4-2

Questions for a person

- **Q-1** (AC-3): Should any order between $50.00 and $74.99 ever ship free, for example through a legacy promotion or for particular customer types? The spec says the old rule is removed, but the diff adds a legacy $50 branch.
- **Q-2** (AC-1): How should subtotals be rounded or compared when they have more than two decimals, such as 74.995? The spec does not define this.
- **Q-3** (new): Is the threshold checked against the subtotal before or after discounts, taxes or promo codes? The spec does not say.
- **Q-4** (new): What is the expected shipping for a zero or negative subtotal, such as an empty cart or a refund? The spec does not say.
- **Q-5** (AC-4): Do all customer types, such as new, guest and returning customers, follow the same threshold? AC-4 only covers returning-customer-01.
- **Q-6** (new): Does the quote API (AC-1 to AC-3) use the same shipping function as the browser receipt path (AC-4)? Is any other threshold copy, such as in the UI or docs, expected to change too?

AI suggestions (not used for labels)

The expected values here are the AI's guesses, not facts. They never change a label or the verdict.

SCENARIO	INPUT	AI GUESS	ACTUAL	STATUS
S-AI-1 AC-3	via=api.quote_order, subtotal=50.00	shipping 5.99	subtotal 50.00, shipping 0.00, discount 0.00, discount_label none, total 50.00	FAIL
S-AI-2 AC-3	via=api.quote_order, subtotal=60.00	shipping 5.99	subtotal 60.00, shipping 0.00, discount 0.00, discount_label none, total 60.00	FAIL
S-AI-3 AC-2	via=api.quote_order, subtotal=74.99	shipping 5.99	subtotal 74.99, shipping 0.00, discount 0.00, discount_label none, total 74.99	FAIL
S-AI-4 AC-1	via=api.quote_order, subtotal=75.00	shipping 0.00	subtotal 75.00, shipping 0.00, discount 0.00, discount_label none, total 75.00	PASS
S-AI-5 AC-3	via=api.quote_order, subtotal=0.00	shipping 5.99	HTTP 422: {"error": {"code":"VALIDATION_ERROR","message":"Request validation failed","details":[{"field":"items.0.unit_price","message":"Input should be greater than 0"}]}}	ERROR

Files

Changed: `app/config.py`, `app/pricing/shipping.py`

Not covered by any criterion: `app/pricing/shipping.py`

In a high-risk path: none

Existing tests: 57/58 passed (context only).
Challenger AI: saved: model unknown, prompt 984d7b042982
Commit `b7d0762` against `main` · run `r-20261104-115730` · 2026-11-04T11:57:57Z
Evidence Gate decides how much human review a change needs. It never removes human review.